

Application resource management

300-level live demo script



Introduction

In this growing digital economy, the application is the business. Application performance, therefore, is one of highest CIO priorities.

Home Robots Inc. is a fully digital company selling innovative “household chores” robots globally via its RobotShop online marketplace. Clients browse and purchase through RobotShop’s microservices-based cloud native app. Promotions and other marketing events, however, generate unpredictable load patterns. This often results in poor application performance and bad customer experiences. IT Ops teams tend to either over- or under-provision resources based on best guesses - which is highly inefficient, costly and risky.

In this demo, I’ll show you how IBM Turbonomic, a solution based on the principles of Application Resource Management (ARM), helps SREs and IT Ops teams proactively assure application performance and operational efficiency across their mission critical deployments. We will:

- See how Turbonomic stitches together a full-stack view of your multi-cloud software deployments
- Examine the resource optimization recommendations generated by Turbonomic
- Demonstrate how to automate the execution of recommended actions

1 - Getting a global view of the applications and their infrastructure dependencies

1.1 - Examine the global supply chain

Narration

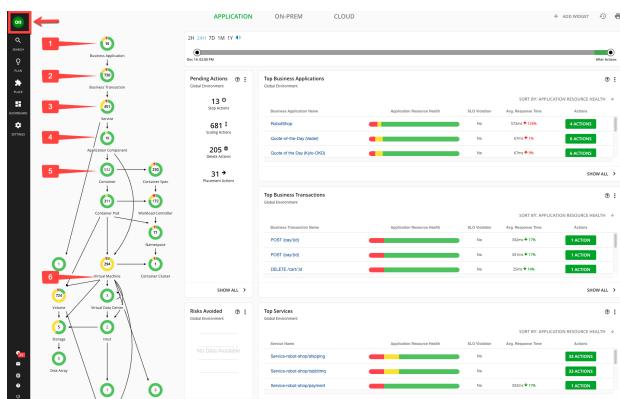
Turbonomic requires configuration and operational performance data to make resource optimization recommendations. The more data the better.

The RobotShop application is being observed by Instana. Turbonomic ingests data from Instana and builds a common data model to stitch together a graphical view of the application and its resource dependencies. In the Turbonomic nomenclature, this is called the “supply chain.”

The global supply chain models the dynamic relationship of the managed application and its dependent infrastructure layers.

Action 1.1.0

- Log in to the Turbonomic instance and click the **On** button in the upper left corner.
- **NOTE:** The next six steps will refer to the graphic below.



Action 1.1.1

- Hover the cursor over the **Business Application** entity (1).

Narration

- The Business Applications entity shows the business applications of which Turbonomic is aware.

Action 1.1.2

- Hover the cursor over the **Business Transaction** entity (2).

Narration

- The Business Transaction entity shows logical business functions that an end-user would execute (such as a purchase or add-to-cart). Business applications are composed of these business transactions.

Action 1.1.3

- Hover the cursor over the **Service** entity (3).

Narration

- A service is basically a REST endpoint, and transactions use the services.

Action 1.1.4

- Hover the cursor over the **Application Component** entity (4).

Narration

- Services are hosted and executed in an application component, such as a Java virtual machine (JVM).

Action 1.1.5

- Hover the cursor over the **Container** entity (5).

Narration

- Application components run on a container platform like Kubernetes.

Action 1.1.6

- Hover the cursor over the **Virtual Machine** entity (6).

Narration

- Application platforms are hosted in virtualized environments like vSphere.

1.2 - Explore the top business applications view

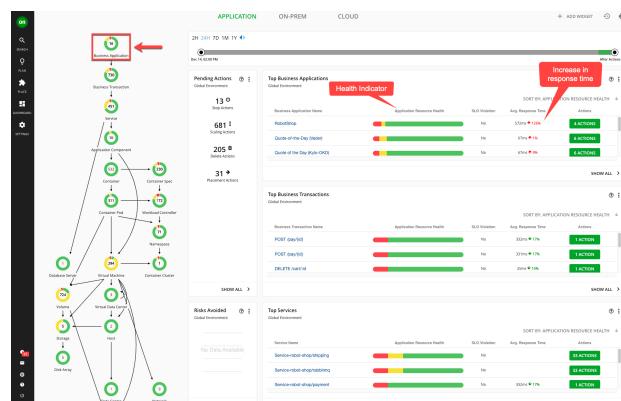
Narration

Turbonomic displays the applications in order of risk.

In each of the entities, such as the Business Application entity, the colors that make up the circle reflect the current health of the entities. “Green” is healthy, “yellow” represents efficiency recommendations, and “red” represents performance recommendations..

Action 1.2.1

- Click the **Business Application** entity.



Narration

We see that there's an increase in the average response time.

The current status indicates that there are some critical performance issues as well as some areas to improve efficiency.

The Actions button enables you to take the recommended actions directly from Turbonomic.

NOTE: We won't click the Actions button at this time.

Action 1.2.2

- Click the X in the upper right corner to close the **Application: Business Application** page.



2 - Drilling into the RobotShop application

2.1 - Examine RobotShop resource dependencies

Narration

Now that we have a broad understanding of the global view, let's examine the health of the RobotShop application. This is called "scoping."

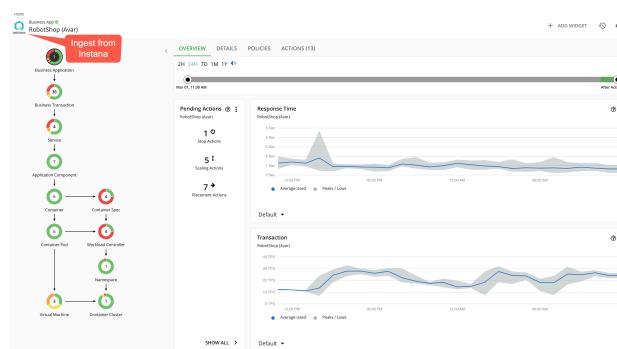
Action 2.1.1

- Click the **RobotShop** link to scope to RobotShop.



Action 2.1.2

- Point out that we are scoped to RobotShop, with data coming from Instana.



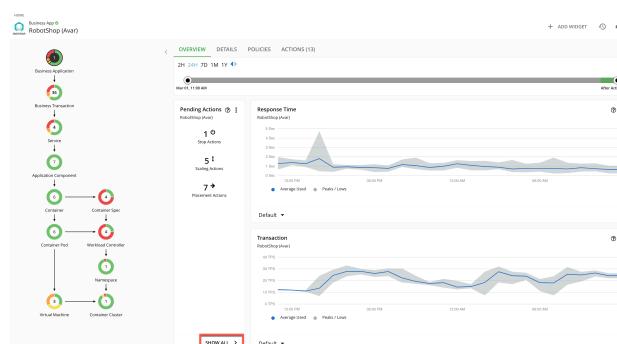
Narration

The supply chain can ingest data from a number of APM (Application Performance Monitoring) tools. In this case, the supply chain is scoped to RobotShop, and the charts provide a quick view of RobotShop's overall operating health.

Since RobotShop is a Kubernetes-based cloud native application, all the entities in the supply chain are specific to a container infrastructure.

Action 2.1.3

- On the **Pending Actions** chart, click **SHOW ALL**.



Narration

The Turbonomic engine performs an ongoing (real-time) holistic analysis of the environment, generating resource optimization recommendations (and associated actions) that you can follow to resolve and avoid emerging problems.

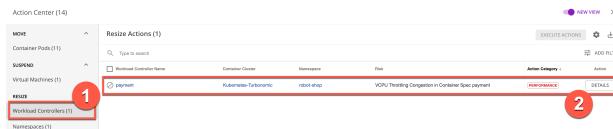
Here we see container resize actions, categorized as ‘performance’ and ‘efficiency’ actions. These are displayed for you to either investigate further or execute manually.

Performance: Container resize *UP* actions are typically performance-centric actions that are driven to resolve an underlying resource congestion issue.

Efficiency: Container resize *DOWN* actions are typically efficiency-centric actions that are pointing to a resource optimization opportunity, likely a consequence of resource over-provisioning.

Action 2.1.4

- On the **Action Center** panel, under **RESIZE**, select **Workload Controllers** (1). Then, click **DETAILS** (2) in the **payment** row.
- NOTE:** If there are multiple **payment** rows, select the first one.
- NOTE:** If the view does not match the screenshot below, click **NEW VIEW** at the upper right corner to enable the new view.



3 - Understanding the resource optimization recommendations

3.1 - Analyze a container right-sizing performance recommendation

Narration

Let's explore one of the performance recommendations in more detail.

Action 3.1.1

- On the **Action Details** page, if the details are not already expanded, click **Expand Details**.

The screenshot shows the 'Action Center (14)' interface. A specific action titled 'Resize VCPU Limit for Workload Controller payment' is selected. The 'IMPACTED CONTAINER SPEC' section is expanded, revealing 'EXPAND DETAILS'. Below it, 'WORKLOAD CONTROLLER DETAILS' show the container name 'payment' and namespace 'robot-shop'. The 'RELATED ACTIONS' section contains a link to 'Resize up VCPU Limit Quota for Namespace robot-shop from 4,000 mCores to 6,000 mCores'.

Action 3.1.2

- Review the performance action details. Point out the recommendation to upsize the virtual CPU limit from 200 mCores to 500 mCores.
- NOTE:** Since this is a dynamic system, the metrics displayed on your screen may not necessarily reflect the exact numbers in the screenshot.

The screenshot shows the 'Action Details' page for the same resize action. It includes a 'VIRTUAL CPU LIMIT PERCENTILE AND AVG. UTILIZATION' chart comparing current vs projected values. The chart shows utilization dropping significantly after March 1st. Below the chart, 'CONTAINER SPEC - IMPACT FROM ALL ACTIONS' lists the current 'Virtual CPU Limit Daily Avg.' at 200 mCores and the 'Projected Virtual CPU Limit 99th Percentile' at 500 mCores. At the bottom, 'WORKLOAD CONTROLLER DETAILS' provide additional context about the container and its cluster.

Narration

CPU throttling directly impacts the response time of a service and therefore user experience. Turbonomic observes and tracks CPU throttling. In the case of RobotShop, the response time latency of the payment service can adversely impact the ecommerce checkout experience. Given the rate of throttling being observed, Turbonomic recommends increasing the CPU limit from 200 mCores to 500 mCores, thereby reducing the CPU throttling from 19.2% to 3.1%. Proactively alleviating CPU congestion pressures by providing these analytics-driven recommendations helps keep the application in the desired state where it meets its defined service-level objectives (SLOs).

Action 3.1.3

- Click the X in the upper right corner to close the **Action Details** page.



3.2 - Analyze a proactive workload redistribution recommendation

Narration

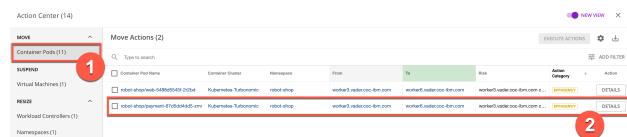
The underlying cloud native environment to which the RobotShop microservices are deployed is highly dynamic.

The initial node to which a pod was placed may not always remain the most optimal place to continue executing this workload. An unhealthy application can cause cascading issues, potentially impacting healthy neighboring pods. Additional node capacity may present alternative optimal placement options that were not available at the time when this pod was initially placed.

Continuously and proactively redistributing workloads, in line with shifting load patterns and available capacity, helps assure application performance and infrastructure operational efficiency.

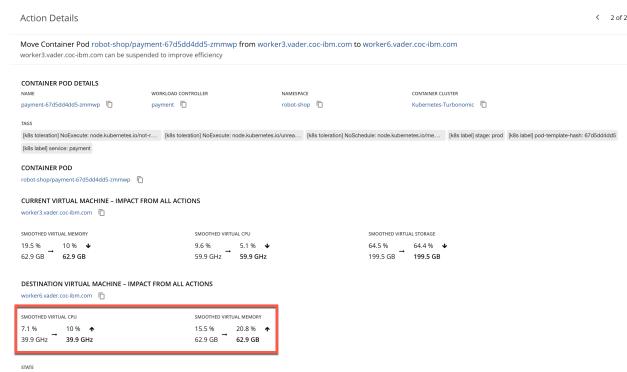
Action 3.2.1

- On the **Action Center** panel, under **MOVE**, select **Container Pods** (1). Then, click **DETAILS** (2) in the **robot-shop/payment** row.
- NOTE:** If there are multiple **robot-shop/payment** rows, select the first one.



Action 3.2.2

- Review the efficiency action details. Point out the recommendation to relocate the payment pod from where it is currently running on worker node 3 to worker node 6.
- NOTE:** Since this is a dynamic system, the metrics displayed on your screen may not necessarily reflect the exact numbers in the screenshot.



Move Container Pod robot-shop/payment-67d5d4d5-zmmwp from worker3.vader.coc.ibm.com to worker6.vader.coc.ibm.com
worker3.vader.coc.ibm.com can be suspended to improve efficiency

STATE	Resource	Current Usage (%)	Impact (%)
idle	CPU	10%	10%
idle	Memory	62.9 GB	62.9 GB
idle	Storage	59.9 GHz	59.9 GHz
idle	Network	59.9 GHz	59.9 GHz

Narration

If the underlying node is getting congested, it will attempt to identify an alternate node with more capacity. If, on the other hand, as is the case here, the underlying node is underutilized, it makes sense to proactively redistribute the workload to other appropriate nodes to improve operational efficiency.

The Turbonomic analysis engine computes the current and possible future state of the cluster if the Move action is accepted for execution. As we can observe here, the migration of workloads from worker node 3 to worker node 6 will result in a marginal increase of virtual CPU utilization from 7.1% to 10% and an increase of virtual memory consumption from 15.5% to 20.8% of worker node 6.

Continuous redistribution of workloads helps better optimize overall cluster resources in terms of performance and cost of ownership.

Action 3.2.3

- Click the X in the upper right corner to close the **Action Details** page.

The screenshot shows the 'Action Details' page with the following details:

- CONTAINER POD DETAILS:**
 - Name: payment-67d5dd4dd5-zmmpw
 - WORKLOAD CONTROLLER: payment
 - NAMESPACE: payment-67d5dd4dd5-zmmpw
 - CONTAINER CLUSTER: Kubernetes/Turbonomic
- INFO:** [Info] [Info] [Info] [Info] [Info] [Info] [Info] [Info] [Info] [Info]
- CURRENT VIRTUAL MACHINE - IMPACT FROM ALL ACTIONS:** worker3.vader.cos-dlm.com

3.3 - Analyze an intelligent cluster scaling recommendation

Narration

The Turbonomic analysis engine is continuously exploring opportunities to optimize overall cluster efficiency, essentially balancing cluster capacity with workload demand.

Pods consume resources from the underlying nodes on which they are placed. Nodes are represented as virtual machines (VMs) in the supply chain. When pods begin experiencing performance issues due to diminishing resources at the underlying node level, Turbonomic will provide early recommendations to provision additional cluster capacity.

On the other hand, it will seek out opportunities to consolidate workloads on a fewer number of nodes, thereby driving down operational costs.

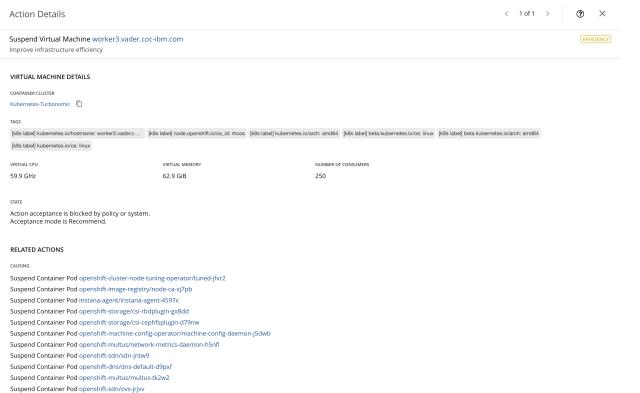
Action 3.3.1

- On the **Action Center** panel, under **SUSPEND**, select **Virtual Machines** (1). Then, click **DETAILS** (2) in the **vader.coc-ibm.com** row.



Action 3.3.2

- Review the action details.



Narration

Executing the Suspend action will result in reclaiming and possibly repurposing the compute resources of worker node 3.

These actions enhance overall cluster operational efficiency and naturally yield significant costs savings, while not compromising application performance and availability.

Action 3.3.3

- Click the X in the upper right corner to close the **Action Details** page.

Action Details

Suspend Virtual Machine worker3.vader.coc-ibm.com

Improve infrastructure efficiency

VIRTUAL MACHINE DETAILS

CPU

Kubernetes Tutorials: 1

VGA

[Jobs label] kubernetes.iohostname: worker3.vader.coc-ibm.com [Jobs label] node.openshift.io/os_id: rhcos_4.14.0 [Jobs label] kubernetes.io/arch: amd64 [Jobs label] beta.kubernetes.io/cos: Intel [Jobs label] kubernetes.io/vcpu: 16

VIRTUAL CPU

59.9 GHz

VIRTUAL MEMORY

62.9 GB

NUMBER OF CONSUMERS

250

4 - Automating actions

4.1 - Automate the execution of actions

Narration

Though Turbonomic allows you to initiate an action natively from the platform with the click of a button, it is a best practice to automate the execution of these actions.

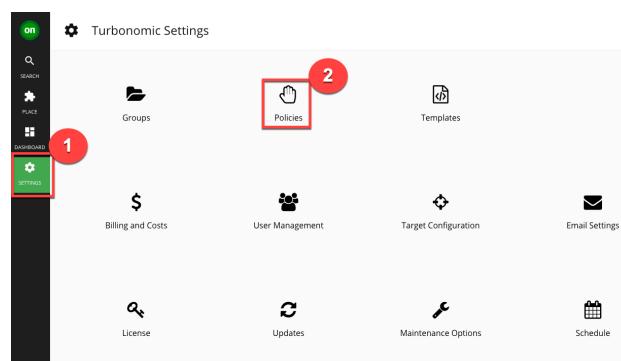
We will now define a policy that enables you to automate the platform-derived actions without having to jump into multiple tools. This significantly enhances operator productivity.

Action 4.1.1

- Log out and log back in using the following credentials:
- Username: **cocadmin**
- Password: **CoC#Rulz!**

Action 4.1.2

- Click **SETTINGS** (1) and select **Policies** (2).



Action 4.1.3

- Click **NEW AUTOMATION POLICY**.
- **NOTE:** If you do not see the **NEW AUTOMATION POLICY** button, it means that the system is currently running in a restricted **read-only** mode. In this case, skip Actions 4.1.4-4.1.6 and proceed to the narration.
- **NOTE:** Actions 4.1.4-4.1.6 represent how an automation policy can be configured and put into effect.



Action 4.1.4

- Select **Container Pod**.

Select policy type

Application Component

Business Application

Business Transaction

Container

Container Pod

Container Spec

Database Server

Disk Array

Host

Service

Storage

Virtual Machine

Volume

Workload Controller

Action 4.1.5

- Give the policy a custom **NAME**. Then, expand **AUTOMATION AND ORCHESTRATION** (2) and click **ADD ACTION** (3) to define how an action is accepted

< Configure Container Pod Policy

The screenshot shows a configuration interface for a 'Container Pod Policy'. At the top, there's a back arrow and the title 'Configure Container Pod Policy'. Below that is a 'NAME' input field containing 'RobotShop Container Pod Automation Policy', which is highlighted with a red box and has a red circle with the number '1' above it. Underneath the name is a 'SCOPE' section with a plus sign and 'ADD CONTAINER POD GROUPS'. The next section is 'POLICY SCHEDULE' with a plus sign and 'AUTOMATION AND ORCHESTRATION'. A red circle with the number '2' is placed over the 'AUTOMATION AND ORCHESTRATION' section, which is also highlighted with a red box. A descriptive text 'Defines how actions are accepted.' is shown below it. At the bottom, there's a red box containing a plus sign and 'ADD ACTION', with a red circle with the number '3' placed over it.

Action 4.1.6

- Fill out the **Automation and Orchestration** panel:
 - Set **ACTION TYPE** to **Move** (1).
 - Set **ACTION GENERATION** to **Generate Actions** (2).
 - Set **ACTION ACCEPTANCE** to **Automatic** (3).
- IMPORTANT NOTE:** Do **NOT** click **Submit**, as this is a read-only environment.

Automation and Orchestration

X

ACTION TYPE

Move 1

ACTION GENERATION

Generate Actions 2

ACTION ACCEPTANCE

Automatic 3

BEFORE EXECUTION

Do Nothing

ACTION EXECUTION

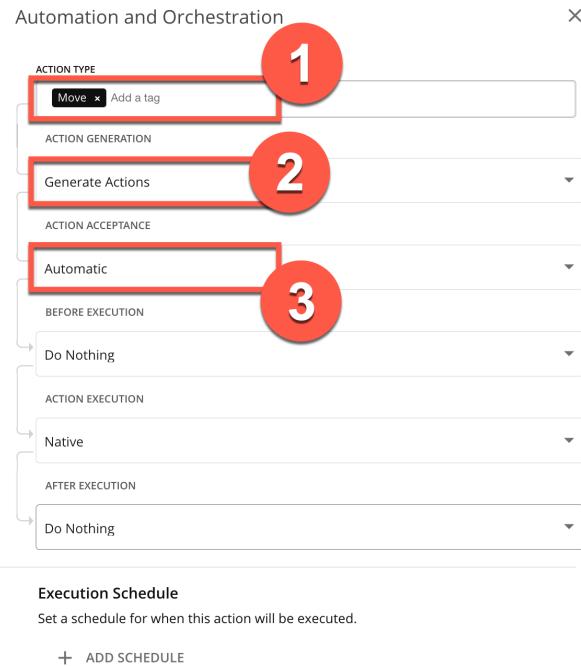
Native

AFTER EXECUTION

Do Nothing

Execution Schedule
Set a schedule for when this action will be executed.

+ ADD SCHEDULE



Narration

Once the automation policy is saved, it will go into effect. All configured actions will now be executed automatically.

The main benefit and best practice of Turbonomic is to execute the platform-derived actions automatically. The underlying goal is to reduce or remove human intervention and leverage automation to maintain application performance and improve operational efficiency.

Although we have demonstrated how actions can be initiated automatically from Turbonomic, typically IT organizations manually execute actions until they are confident that the actions are trustworthy. As the organization's level of comfort matures over time, actions are executed in an increasingly automated fashion.

Summary

In this demo, we showed how Home Robots Inc. leveraged Turbonomic to assure application performance and improve the operational efficiency of the supporting infrastructure.

We demonstrated how Turbonomic can augment a container platform and provide additional high-value capabilities, such as full-stack visibility, analytics-driven resource optimization recommendations, and automation of executing the recommended actions.

Thank you for attending today's presentation.